

Adrian Wedd

AI Safety Engineer & Independent Researcher

Cygnet, Tasmania, Australia

[GitHub](#) [LinkedIn](#) [Email](#) [PDF](#) [Print](#) [Watch Me Work](#)

ABOUT

AI Safety Engineer and Independent Researcher. Three years empirical research on frontier AI models, focused on red-teaming, evaluation frameworks, and failure-first methodology. Seven years translating complex technical findings into actionable insights for government decision-makers.

CORE COMPETENCIES

AI Safety & Evaluation

Red-teaming, adversarial testing, failure-first methodology

Frontier AI Models

Claude API, multi-agent systems, evaluation frameworks

Risk Assessment

Pre-mortem analysis, FMEA, failure mode identification

Policy Translation

Technical findings to actionable governance for decision-makers

EXPERIENCE

Applications Specialist / Acting Senior Change Analyst

Homes Tasmania 2018 - Present

Complex socio-technical systems analysis for Tasmania's housing and community services portfolio. Developed Homes Tasmania's first Generative AI usage policy in 2023.

- Developed Homes Tasmania's GenAI usage policy, procedure, and training materials
- Led cybersecurity initiatives improving system security
- Systems integration and API development using RESTful APIs and Python

Python, PowerShell, JavaScript, RESTful APIs, Azure

PROJECTS

ADHDo

[GitHub](#)

AI safety system with Claude integration, confidence gating, and multi-tool orchestration.

Python Claude API Safety Systems

Agentic Research Engine

[GitHub](#)

Multi-agent evaluation framework with LangGraph orchestration and self-correction loops.

Python LangGraph Multi-Agent Systems

SKILLS

Programming Languages

Python Primary

JavaScript / TypeScript Primary

AI & Safety

Frontier AI Models Primary

Red-Teaming Primary

Evaluation Frameworks Primary

Infrastructure

Systems Integration Primary

Cybersecurity Primary

Research Methods

Risk Assessment Primary

Technical Writing Primary

ACHIEVEMENTS

Failure-First Research Program

Developed comprehensive failure-first evaluation methodology for agentic AI systems.

2022-Present

Published Research: Organisational AI Governance

Published research on structural barriers to acting on AI safety evidence in organisations.

2025

AI Governance Policy Development

Developed Homes Tasmania's first Generative AI usage policy and training materials.

2023

Adrian Wedd

AI Safety Engineer & Independent Researcher

Last updated: Feb 20, 2026, 10:12 AM

© 2026 Adrian Wedd